

# Day 4: Classification

## Summer STEM: Machine Learning

Department of Electrical and Computer Engineering  
NYU Tandon School of Engineering  
Brooklyn, New York

July 15, 2021

# Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Logistic Regression
- 4 Lab: Diagnosing Breast Cancer
- 5 Multiclass Classification
- 6 Lab: Iris Dataset

# Review

- Machine learning pipeline:
  - Process Data
  - Train on training data
  - Test on testing data
- Is it possible have a high accuracy for the training data and a low accuracy for the testing data? What should we do?

# Review

- Imagine you are preparing for the SATs and you come across a book full of practice questions you did not understand how to solve any of the problems. However, you memorized all of the answers.
- What do you think will happen if you try to solve practice questions in a different book.
- Why are you studying actual problem solving techniques instead of just memorizing solutions from practice questions?
- Assuming you have an eidetic memory will memorizing solutions from practice questions be a good strategy?

# Review

- $J(w) = \frac{1}{N} \|Y - Xw\|^2 + \lambda \|w\|^2$
- $w = [10000, 20000, 30000, 10000]$  does this look good?

# Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Logistic Regression
- 4 Lab: Diagnosing Breast Cancer
- 5 Multiclass Classification
- 6 Lab: Iris Dataset

# Motivation

- Cannot rely on closed form solutions
  - Computation efficiency: operations like inverting a matrix is not efficient
  - For more complex problems such as neural networks, a closed-form solution is not always available
- Need an optimization technique to find an optimal solution
  - Machine learning practitioners use **gradient**-based methods

# Gradient Descent Algorithm

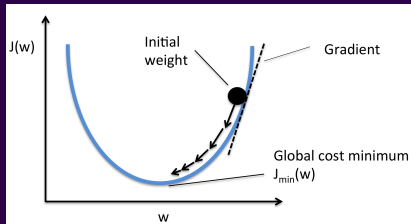
## ■ Update Rule

*Repeat*{

$$\mathbf{w}_{new} = \mathbf{w} - \alpha \nabla J(\mathbf{w})$$

}

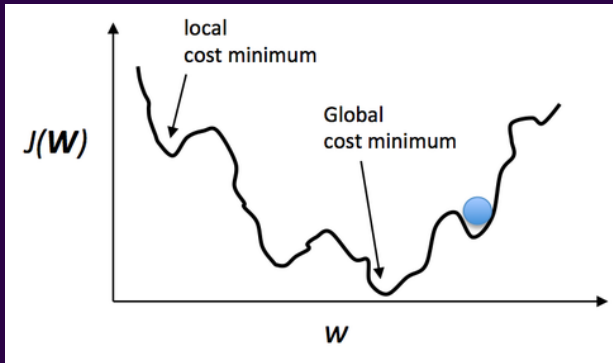
$\alpha$  is the learning rate



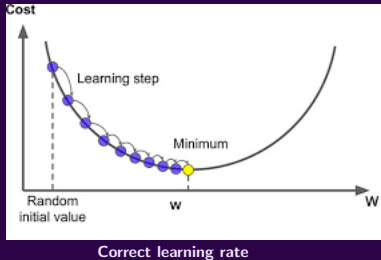
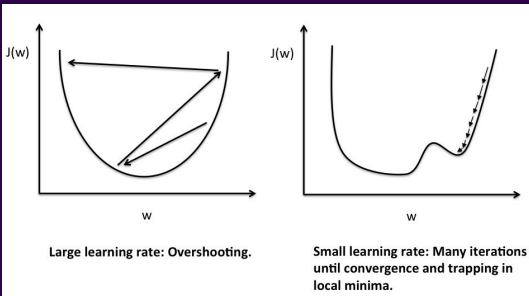


# General Loss Function Contours

- Most loss function contours are not perfectly parabolic
- Our goal is to find a solution that is very close to global minimum by the right choice of hyper-parameters



# Understanding Learning Rate



# Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Logistic Regression**
- 4 Lab: Diagnosing Breast Cancer
- 5 Multiclass Classification
- 6 Lab: Iris Dataset

# Classification Vs. Regression

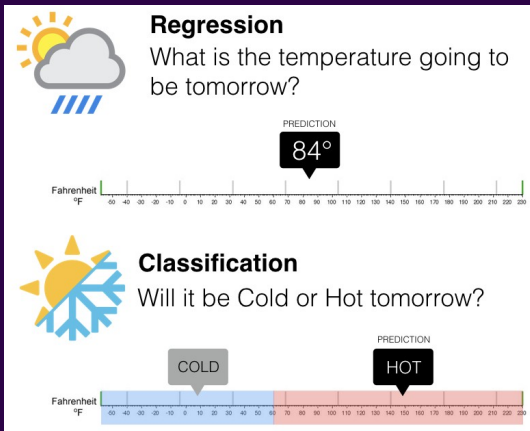

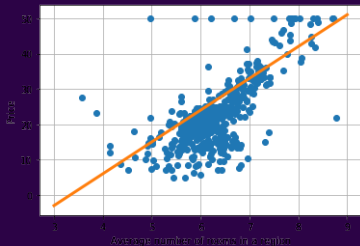


Figure: <https://www.pinterest.com/pin/672232681855858622/?lp=true>

# Classification Vs. Regression



(a) Breast cancer dataset

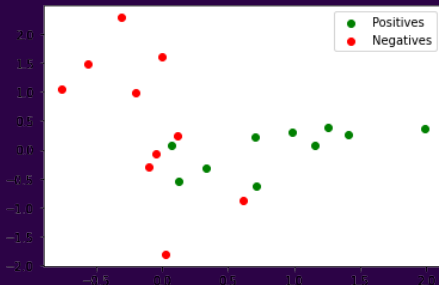


(b) Boston Housing dataset

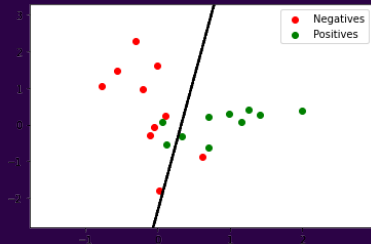
# Classification

Given the dataset  $(x_i, y_i)$  for  $i = 1, 2, \dots, N$ , find a function  $f(x)$  (model) so that it can predict the label  $\hat{y}$  for some input  $x$ , even if it is not in the dataset, i.e.  $\hat{y} = f(x)$ .

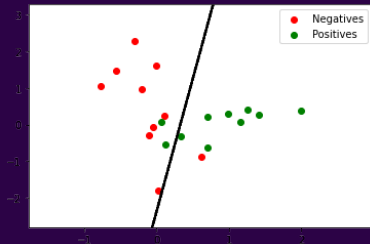
- Positive :  $y = 1$
- Negative :  $y = 0$



# Decision Boundary



# Decision Boundary

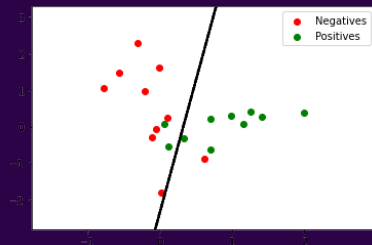


- Evaluation metric :

$$\text{Accuracy} = \frac{\text{Number of correct prediction}}{\text{Total number of prediction}}$$

- What is the accuracy in this example ?





■ Evaluation metric :

$$\text{Accuracy} = \frac{\text{Number of correct prediction}}{\text{Total number of prediction}} = \frac{15}{20} = 0.75 = 75\%$$

# Need for a new model

- What would happen if we used the linear regression model :

$$\hat{y} = w_0 + w_1x$$

# Need for a new model

- What would happen if we used the linear regression model :

$$\hat{y} = w_0 + w_1x$$

- $y$  is 0 or 1
- $\hat{y}$  will take any value between  $-\infty$  and  $\infty$

# Need for a new model

- What would happen if we used the linear regression model :

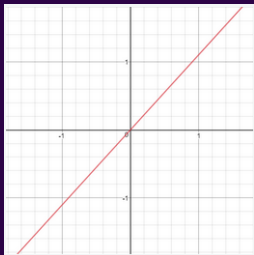
$$\hat{y} = w_0 + w_1 x$$

- $y$  is 0 or 1
- $\hat{y}$  will take any value between  $-\infty$  and  $\infty$
- It will be hard to find  $w_0$  and  $w_1$  that make the prediction  $\hat{y}$  match the label  $y$ .


# Sigmoid Function

- By applying the sigmoid function, we enforce  $0 \leq \hat{y} \leq 1$

$$\hat{y} = \text{sigmoid}(w_0 + w_1x) = \frac{1}{1 + e^{-(w_0 + w_1x)}}$$



(a) Linear model



(b) Sigmoid model

# A new loss function

- Binary cross entropy loss :

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N \left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right]$$

pause

- What happens if  $y_i = 0$  :

$$\left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right] = ?$$

# A new loss function

- Binary cross entropy loss :

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N \left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right]$$

- If  $y_i = 0$  :

$$\left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right] = -\log(1 - \hat{y}_i)$$

# A new loss function

- Binary cross entropy loss :

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N \left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right]$$

- If  $y_i = 0$  :

$$\left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right] = -\log(1 - \hat{y}_i)$$

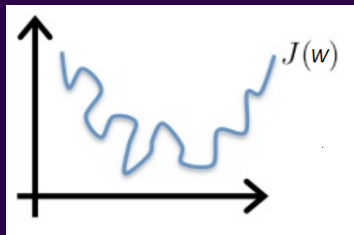
- If  $y_i = 1$  :

$$\left[ -y_i \log(\hat{y}_i) - (1 - y_i) \log(1 - \hat{y}_i) \right] = -\log(\hat{y}_i)$$

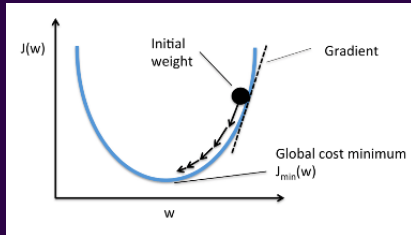


# MSE vs Binary cross entropy loss

- MSE of a logistic function has many local minima.
- The Binary cross entropy loss has only one minimum.



(a) MSE



(b) Binary cross entropy loss

# Classifier

$$\hat{y} = \text{sigmoid}(w_0 + w_1 x) = \frac{1}{1 + e^{-(w_0 + w_1 x)}}$$

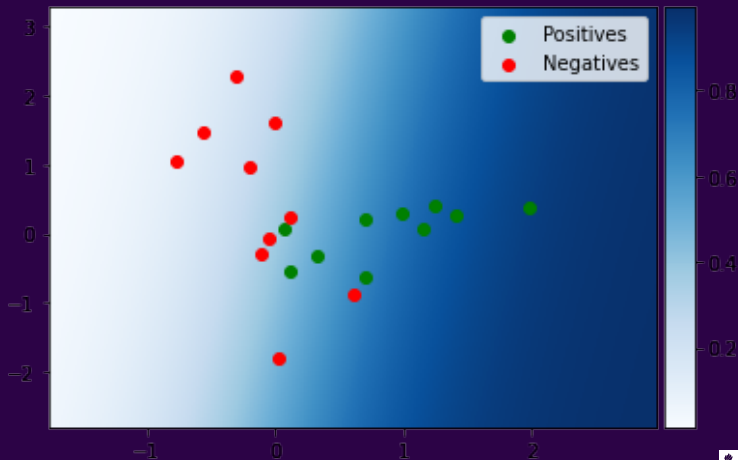
- How to deal with uncertainty ?
  - Thanks to the sigmoid,  $\hat{y} = f(x)$  is between 0 and 1.

# Classifier

$$\hat{y} = \text{sigmoid}(w_0 + w_1 x) = \frac{1}{1 + e^{-(w_0 + w_1 x)}}$$

- How to deal with uncertainty ?
  - Thanks to the sigmoid,  $\hat{y} = f(x)$  is between 0 and 1.
- If  $\hat{y}$  is close to 0, the data is probably negative
- If  $\hat{y}$  is close to 1, the data is probably positive
- If  $\hat{y}$  is around 0.5, we are not sure.

# Classifier



# Decision Boundary

- Once, we have a classifier outputting a score  $0 < \hat{y} < 1$ , we need to create a decision rule.

# Decision Boundary

- Once, we have a classifier outputting a score  $0 < \hat{y} < 1$ , we need to create a decision rule.
- Let  $0 < t < 1$  be a Threshold :
  - If  $\hat{y} > t$ ,  $\hat{y}$  is classified as positive.
  - If  $\hat{y} < t$ ,  $\hat{y}$  is classified as negative.

# Decision Boundary

- Once, we have a classifier outputting a score  $0 < \hat{y} < 1$ , we need to create a decision rule.
- Let  $0 < t < 1$  be a Threshold :
  - If  $\hat{y} > t$ ,  $\hat{y}$  is classified as positive.
  - If  $\hat{y} < t$ ,  $\hat{y}$  is classified as negative.
- How to choose  $t$  ?

# Impact of the threshold

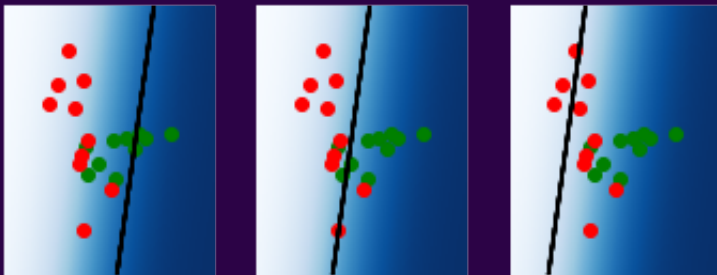


Figure:  $t = 0.2, 0.5, 0.8$



# Performance metrics for a classifier

- Accuracy of a classifier: percentage of correct classification
- Why accuracy alone is not a good measure for assessing the model ?


# Performance metrics for a classifier

- Accuracy of a classifier: percentage of correct classification
- Why accuracy alone is not a good measure for assessing the model ?
  - Example: A rare disease occurs 1 in ten thousand people
  - A test that classifies everyone as free of the disease can achieve 99.999% accuracy when tested with people drawn randomly from the entire population

# Types of Errors in Classification

- Correct predictions:
  - True Positive (TP) : Predict  $\hat{y} = 1$  when  $y = 1$
  - True Negative (TN) : Predict  $\hat{y} = 0$  when  $y = 0$
- Two types of errors:
  - False Positive/ False Alarm (FP):  $\hat{y} = 1$  when  $y = 0$
  - False Negative/ Missed Detection (FN):  $\hat{y} = 0$  when  $y = 1$

# Example



- How many True Positive (TP) are there ?
- How many True Negative (TN) are there ?
- How many False Positive (FP) are there ?
- How many False Negative (FN) are there ?

# Other metrics

- Sensitivity/Recall/TPR (How many positives are detected among all positive?)

$$\frac{TP}{TP + FN}$$

- Precision (How many detected positives are actually positive?)

$$\frac{TP}{TP + FP}$$

# Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Logistic Regression
- 4 Lab: Diagnosing Breast Cancer
- 5 Multiclass Classification
- 6 Lab: Iris Dataset

# Lab: Diagnosing Breast Cancer

- We're going to use the breast cancer dataset to predict whether the patients' scans show a malignant tumour or a benign tumour.
- Let's try to find the best linear classifier using logistic regression.

# Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Logistic Regression
- 4 Lab: Diagnosing Breast Cancer
- 5 Multiclass Classification**
- 6 Lab: Iris Dataset



# Multiclass Classification

- Previous model:  $f(\mathbf{x}) = \sigma(\phi(\mathbf{x})w)$
- Representing Multiple Classes:
  - One-hot / 1-of-K vectors, ex : 4 Class
  - Class 1 :  $\mathbf{y} = [1, 0, 0, 0]$
  - Class 2 :  $\mathbf{y} = [0, 1, 0, 0]$
  - Class 3 :  $\mathbf{y} = [0, 0, 1, 0]$
  - Class 4 :  $\mathbf{y} = [0, 0, 0, 1]$

# Multiclass Classification

- Previous model:  $f(\mathbf{x}) = \sigma(\phi(\mathbf{x})w)$
- Representing Multiple Classes:
  - One-hot / 1-of-K vectors, ex : 4 Class
  - Class 1 :  $\mathbf{y} = [1, 0, 0, 0]$
  - Class 2 :  $\mathbf{y} = [0, 1, 0, 0]$
  - Class 3 :  $\mathbf{y} = [0, 0, 1, 0]$
  - Class 4 :  $\mathbf{y} = [0, 0, 0, 1]$
- Multiple outputs:  $f(\mathbf{x}) = \text{softmax}(\phi(\mathbf{x})W)$
- Shape of  $\phi(\mathbf{x})W$  :  $(N, K) = (N, D) \times (D, K)$
- $\text{softmax}(\mathbf{z})_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$

# Multiclass Classification

- Multiple outputs:  $f(\mathbf{x}) = \text{softmax}(\mathbf{z})$  with  $\mathbf{z} = \phi(\mathbf{x})W$

- $\text{softmax}(\mathbf{z})_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$

- Softmax example: If  $\mathbf{z} = \begin{bmatrix} -1 \\ 2 \\ 1 \\ -4 \end{bmatrix}$  then,

$$\text{softmax}(\mathbf{z}) = \begin{bmatrix} \frac{e^{-1}}{e^{-1}+e^2+e^1+e^{-4}} \\ \frac{e^2}{e^{-1}+e^2+e^1+e^{-4}} \\ \frac{e^1}{e^{-1}+e^2+e^1+e^{-4}} \\ \frac{e^{-4}}{e^{-1}+e^2+e^1+e^{-4}} \end{bmatrix} \approx \begin{bmatrix} 0.035 \\ 0.704 \\ 0.259 \\ 0.002 \end{bmatrix}$$

# Cross-entropy

- Multiple outputs:  $\hat{\mathbf{y}}_i = \text{softmax}(\phi(\mathbf{x}_i)W)$
- Cross-Entropy:  $J(W) = - \sum_{i=1}^N \sum_{k=1}^K \mathbf{y}_{ik} \log(\hat{\mathbf{y}}_{ik})$
- Example :  $K = 4$

$$\text{If, } \mathbf{y}_i = [0, 0, 1, 0] \text{ then, } \sum_{k=1}^K \mathbf{y}_{ik} \log(\hat{\mathbf{y}}_{ik}) = \log(\hat{\mathbf{y}}_{i3})$$

# Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Logistic Regression
- 4 Lab: Diagnosing Breast Cancer
- 5 Multiclass Classification
- 6 Lab: Iris Dataset

# Lab: Iris Dataset

- Open `demo_iris.ipynb`