

Day 6: Neural Networks

Summer STEM: Machine Learning

Department of Electrical and Computer Engineering
NYU Tandon School of Engineering
Brooklyn, New York

Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Neural Networks
- 4 Stochastic Gradient Descent
- 5 Overparameterization
- 6 Lab

Machine Learning Problem Pipeline

- 1 Formulate the problem: regression, classification, or others?
- 2 Gather and visualize the data
- 3 Design the model and the loss function
- 4 Train your model
 - (a) Perform feature engineering
 - (b) Construct the design matrix
 - (c) Choose regularization techniques
 - (d) Tune hyper-parameters using a validation set
 - (e) If the performance is not satisfactory, go back to step (a).
- 5 Evaluate the model on a test set

Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Neural Networks
- 4 Stochastic Gradient Descent
- 5 Overparameterization
- 6 Lab

Motivation

- Cannot rely on closed form solutions
 - Computation efficiency: operations like inverting a matrix is not efficient
 - For more complex problems such as neural networks, a closed-form solution is not always available
- Need an optimization technique to find an optimal solution
 - Machine learning practitioners use **gradient**-based methods

Gradient Descent Algorithm

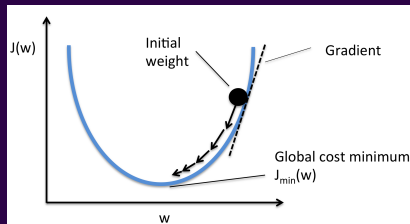
■ Update Rule

Repeat{

$$\mathbf{w}_{new} = \mathbf{w} - \alpha \nabla J(\mathbf{w})$$

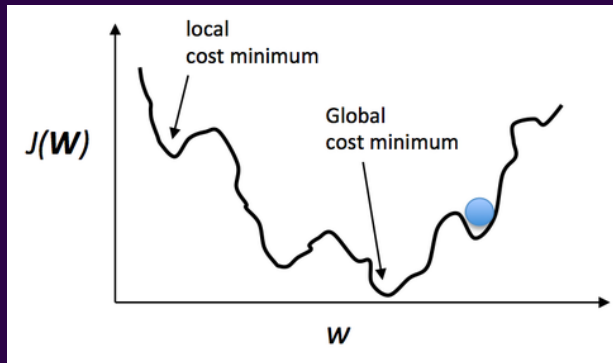
}

α is the learning rate

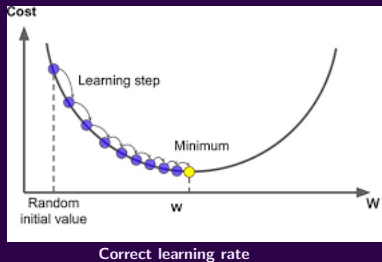
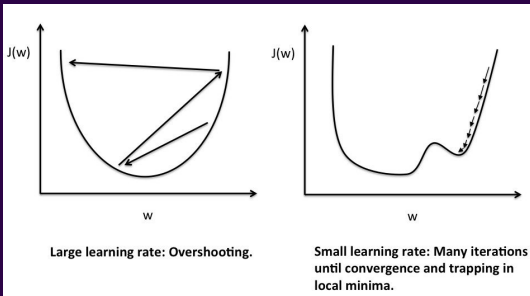


General Loss Function Contours

- Most loss function contours are not perfectly parabolic
- Our goal is to find a solution that is very close to global minimum by the right choice of hyper-parameters




Understanding Learning Rate



Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Neural Networks**
- 4 Stochastic Gradient Descent
- 5 Overparameterization
- 6 Lab

Biological Neuron




Source: David Baillot/UC San Diego

Mathematical Neuron: Perceptron

Biological Neuron:

- A neuron can receive electrochemical signals from other neurons;
- A neuron fires once its accumulated electric charge passes a certain threshold.
- Neurons that fire together wire together.

Mathematical Neuron (Perceptron)



Relation to Logistic Regression

What if we use the sigmoid function as the activation?

$$f(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$$


Decision boundary is a line: how is this supposed to revolutionize machine learning?

Multi-layer Perceptron (MLP)

We need more neurons and we need to connect them together!

- Many ways to do that...
- Today: multi-layer perceptron/fully connected feed-forward network.

MLP Example



- What is the shape of the input and output?
- How many parameters does this model have?
- What activation function would you use for the output layer? Why?

More about MLPs

- Many choices for the activation function: Sigmoid, Tanh, ReLU, Swish, etc.
- Many choices for the number of hidden layers and the number of neurons per layer.
- MLPs can approximate any continuous function given enough data.
- MLPs can overfit, but we know many effective ways of regularization.

Exercise: TensorFlow Keras Basics

Open `demo_tf_keras_basics.ipynb`

Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Neural Networks
- 4 Stochastic Gradient Descent
- 5 Overparameterization
- 6 Lab

Deep Learning

What does deep learning stand for?

- Deep: Neural network architectures with many hidden layers.
- Learning: Optimizing model parameters given a dataset.

In general, the deeper the model is, the more parameters we need to learn and the more data is needed.

Large-Scale Machine Learning

For deep learning systems to perform well, large datasets are required

- COCO 330K images
- ImageNet 14 million images

Challenges:

- Memory limitation: GeForce RTX 2080 Ti has 11 GB memory, while ImageNet is about 300 GB.
- Computation: Calculating gradients for the whole dataset is computationally expensive (slow), and we need to do this many times.

Stochastic Gradient Descent

Idea: Instead of calculating the gradients from the whole dataset, do it only on a subset.

- Randomly select B samples from the dataset
- The loss for this subset

$$\tilde{J}(\mathbf{w}) = \frac{1}{B} \sum_{i=1}^B \|y - \hat{y}_i\|^2$$

- Update Rule

Repeat{

$$\mathbf{w}_{new} = \mathbf{w} - \alpha \nabla \tilde{J}(\mathbf{w})$$

}

Yet Another Hyper-Parameter

This gives a noisy gradient

$$\nabla \tilde{J}(\mathbf{w}) = \nabla J(\mathbf{w}) + \epsilon$$

- **SGD**: $B = 1$, gives very noisy gradients
- **(batch) GD**: $B = N$, $\epsilon = 0$, expensive to compute
- **Mini-batch GD**: Pick a small B , typical values are 32, 64, rarely more than 128 for image inputs

Some Noise Helps

Even if we can, we rarely set $B = N$. In fact, some noise in the gradients might help to

- escape from local minima,
- escape from saddle points, and
- improve generalization.

Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Neural Networks
- 4 Stochastic Gradient Descent
- 5 Overparameterization**
- 6 Lab

Overparameterized Models

- Modern deep learning models are heavily overparameterized, i.e. the number of learnable parameters is much larger than the number of the training samples.

Overparameterized Models

- Modern deep learning models are heavily overparameterized, i.e. the number of learnable parameters is much larger than the number of the training samples.
 - ResNet: State-of-the-art vision model, 10-60 million parameters

Overparameterized Models

- Modern deep learning models are heavily overparameterized, i.e. the number of learnable parameters is much larger than the number of the training samples.
 - ResNet: State-of-the-art vision model, 10-60 million parameters
 - GPT-3: State-of-the-art language model, 175 billion parameters


Overparameterized Models

- Modern deep learning models are heavily overparameterized, i.e. the number of learnable parameters is much larger than the number of the training samples.
 - ResNet: State-of-the-art vision model, 10-60 million parameters
 - GPT-3: State-of-the-art language model, 175 billion parameters
- Conventional wisdom: Such models overfit.

Overparameterized Models

- Modern deep learning models are heavily overparameterized, i.e. the number of learnable parameters is much larger than the number of the training samples.
 - ResNet: State-of-the-art vision model, 10-60 million parameters
 - GPT-3: State-of-the-art language model, 175 billion parameters
- Conventional wisdom: Such models overfit.
- It is not the case in practice!

Double Descent Curve



Source: OpenAI

Outline

- 1 Review
- 2 Non-linear Optimization
- 3 Neural Networks
- 4 Stochastic Gradient Descent
- 5 Overparameterization
- 6 Lab**

Let's solve the mini-project with MLPs!

Open `lab_mlp_fish_market_keras.ipynb`